

MNIST Handwritten Digit Classification using Deep Learning (Neural Network)

Importing the Dependencies

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import cv2
#from google.colab.patches import cv2_imshow
from PIL import Image
import tensorflow as tf
tf.random.set_seed(3)
from tensorflow import keras
from keras.datasets import mnist
from tensorflow.math import confusion_matrix
```

Loading the MNIST data from keras.datasets

```
In [3]: (X_train, Y_train), (X_test, Y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)
11490434/11490434 [=====] - 3s 0us/step

```
In [4]: type(X_train)
```

```
Out[4]: numpy.ndarray
```

```
In [6]: # shape of the numpy arrays
print(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape)

(60000, 28, 28) (60000,) (10000, 28, 28) (10000,)
```

Training data = 60,000 Images

Test data = 10,000 Images

Image dimension --> 28 x 28

Grayscale Image --> 1 channel

```
print(X_train[9])
```

[illegible]


```
In [ ]: print(X_train[10].shape)
```

```
(28, 28)
```

```
In [ ]: # displaying the image
```

```
plt.imshow(X_train[25])  
plt.show()
```

```
# print the corresponding label  
print(Y_train[25])
```




2

```
In [9]: # displaying the image
```

```
plt.imshow(X_train[20])  
plt.show()
```

```
# print the corresponding label  
print(Y_train[20])
```



4

Image Labels

```
In [13]: print(Y_train.shape, Y_test.shape)
```

```
(60000,) (10000,)
```

```
In [12]: # unique values in Y_train
print(np.unique(Y_train))
```

```
# unique values in Y_test
print(np.unique(Y_test))
```

```
[0 1 2 3 4 5 6 7 8 9]
```

```
[0 1 2 3 4 5 6 7 8 9]
```

We can use these labels as such or we can also apply One Hot Encoding

All the images have the same dimensions in this dataset, If not, we have to resize all the images to a common dimension

```
In [10]: # scaling the values
```

```
X_train = X_train/255
```

```
X_test = X_test/255
```

```
In [11]: # printing the 10th image
```

```
print(X_train[10])
```

```
[[0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.]
 [0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.]
 [0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.]
 [0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.
  0.  0.  0.  0.  0.  0.]
  ~  ~  ~  ~  ~  ~]
```

Building the Neural Network

In [20]: *# setting up the layers of the Neural Network*

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28,28)),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='sigmoid')
])
```

In [21]: *# compiling the Neural Network*

```
model.compile(optimizer='adam',
              loss = 'sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

In [22]: *# training the Neural Network*

```
model.fit(X_train, Y_train, epochs=10)
```

```
Epoch 1/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.2612 - accuracy:
0.9233
Epoch 2/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.1177 - accuracy:
0.9639
Epoch 3/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0835 - accuracy:
0.9737
Epoch 4/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0656 - accuracy:
0.9792
Epoch 5/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0541 - accuracy:
0.9830
Epoch 6/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0445 - accuracy:
0.9851
Epoch 7/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0374 - accuracy:
0.9875
Epoch 8/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0323 - accuracy:
0.9891
Epoch 9/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0306 - accuracy:
0.9900
Epoch 10/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0252 - accuracy:
0.9917
```

Out[22]: <keras.src.callbacks.History at 0x25a6df89490>

Training data accuracy = 99.17%

Accuracy on Test data:

In [23]: `loss, accuracy = model.evaluate(X_test, Y_test)`
`print(accuracy)`


```
313/313 [=====] - 1s 2ms/step - loss: 0.1005 - accuracy: 0.
9757
0.9757000207901001
```

Test data accuracy = 97.57%

```
In [24]: print(X_test.shape)
```

```
(10000, 28, 28)
```


```
In [28]: plt.imshow(X_test[5])  
plt.show()
```



```
In [30]: print(Y_test[5])
```

```
1
```

```
In [25]: # first data point in X_test
plt.imshow(X_test[0])
plt.show()
```



```
In [32]: print(Y_test[0])
```

7

```
In [33]: Y_pred = model.predict(X_test)
```

313/313 [=====] - 1s 1ms/step

```
In [34]: print(Y_pred.shape)
```

(10000, 10)

```
In [35]: print(Y_pred[0])
```

[2.0773856e-05 3.2481041e-03 4.6754088e-02 7.4509448e-01 1.0255403e-06
7.3025115e-03 1.1680031e-09 9.9999982e-01 2.8102559e-03 4.0035740e-02]

model.predict() gives the prediction probability of each class for that data point

```
In [ ]: # converting the prediction probabilities to class label
```

```
label_for_first_test_image = np.argmax(Y_pred[0])
print(label_for_first_test_image)
```

7

```
In [ ]:
```

```
In [ ]: # converting the prediction probabilities to class label for all test data points
Y_pred_labels = [np.argmax(i) for i in Y_pred]
print(Y_pred_labels)
```

```
5, 7, 1, 0, 1, 0, 5, 0, 2, 5, 4, 2, 5, 4, 4, 0, 0, 0, 2, 0, 1, 2, 5, 4, 5, 0, 7,
8, 9, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 0, 1, 2, 8, 4, 5, 6, 7, 8, 9, 8, 6, 5, 0, 6,
8, 9, 4, 1, 9, 3, 8, 0, 4, 8, 9, 1, 4, 0, 5, 5, 2, 1, 5, 4, 0, 7, 6, 0, 1, 7, 0,
6, 8, 9, 5, 1, 7, 9, 8, 6, 0, 8, 1, 7, 7, 1, 3, 2, 9, 1, 4, 2, 0, 0, 7, 8, 4, 6,
4, 9, 9, 8, 4, 7, 6, 5, 6, 5, 6, 9, 6, 3, 0, 2, 4, 6, 9, 0, 2, 5, 5, 1, 9, 3, 9,
7, 8, 7, 2, 2, 5, 7, 9, 8, 2, 1, 9, 1, 3, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 0, 1, 2,
3, 4, 5, 6, 7, 8, 9, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 1, 2, 6, 5, 3, 0, 7, 0, 4, 1,
4, 3, 6, 7, 2, 3, 1, 2, 1, 2, 9, 6, 0, 1, 3, 0, 2, 7, 5, 7, 6, 2, 9, 1, 9, 0, 6,
0, 6, 0, 2, 0, 6, 1, 5, 8, 4, 3, 0, 1, 5, 4, 4, 8, 5, 7, 5, 7, 8, 3, 4, 8, 8, 5,
2, 9, 7, 1, 3, 8, 1, 0, 7, 5, 3, 6, 9, 4, 7, 7, 9, 9, 3, 4, 4, 3, 8, 6, 2, 0, 1,
2, 3, 4, 5, 6, 7, 8, 9, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 0, 1, 2, 3, 4, 5, 6, 7, 8,
9, 0, 8, 3, 9, 5, 5, 2, 6, 8, 4, 9, 1, 7, 1, 2, 3, 5, 9, 6, 9, 1, 1, 1, 2, 9, 5,
6, 8, 1, 2, 0, 7, 7, 5, 8, 2, 9, 8, 9, 0, 4, 6, 7, 1, 3, 4, 5, 6, 0, 3, 6, 8, 7,
0, 4, 2, 7, 4, 7, 5, 4, 3, 4, 2, 8, 1, 5, 1, 2, 0, 2, 5, 6, 4, 3, 0, 0, 0, 3, 3,
5, 7, 0, 6, 4, 8, 8, 6, 3, 4, 6, 9, 9, 8, 2, 7, 7, 1, 0, 1, 2, 3, 4, 5, 6, 7, 8,
9, 0, 1, 2, 3, 4, 5, 6, 7, 8, 0, 1, 2, 3, 4, 5, 6, 7, 8, 2, 1, 7, 2, 5, 0, 8, 0,
2, 7, 8, 8, 3, 6, 0, 2, 7, 6, 6, 1, 2, 8, 8, 7, 7, 4, 7, 7, 3, 7, 4, 5, 4, 3, 3,
8, 4, 1, 1, 9, 7, 4, 3, 7, 3, 3, 0, 2, 5, 5, 6, 6, 3, 5, 2, 5, 9, 9, 8, 4, 1, 0,
6, 0, 9, 6, 8, 8, 5, 6, 1, 1, 9, 8, 9, 2, 3, 5, 5, 9, 4, 2, 1, 9, 2, 9, 2, 0, 6,
0, 4, 0, 0, 1, 2, 3, 4, 7, 8, 9, 0, 1, 2, 3, 7, 8, 9, 0, 1, 2, 3, 4, 7, 8, 9, 7,
```

Y_test --> True labels

Y_pred_labels --> Predicted Labels

Confusin Matrix

```
In [ ]: conf_mat = confusion_matrix(Y_test, Y_pred_labels)
```


```
In [ ]: print(conf_mat)
```

```
tf.Tensor(
[[ 969    1    1    0    1    1    2    2    3    0]
 [   0 1122    1    2    0    4    2    0    4    0]
 [   5    0  999    1    7    1    2    6   11    0]
 [   0    0    5  966    0   23    0    3    7    6]
 [   0    0    2    0  968    1    4    2    1    4]
 [   2    0    0    4    2  878    4    1    1    0]
 [   6    1    0    0    8   11  929    0    3    0]
 [   1    5    9    0    5    0    0  994    8    6]
 [   2    0    3    6    6   10    2    2  943    0]
 [   3    4    0    8   24   10    1    7   10  942]], shape=(10, 10), dtype=int32)
```



```
In [ ]: plt.figure(figsize=(15,7))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap='Blues')
plt.ylabel('True Labels')
plt.xlabel('Predicted Labels')
```

Out[30]: Text(0.5, 42.0, 'Predicted Labels')



Building a Predictive System

Prediction image link:

<https://camo.githubusercontent.com/3d9666a8f0c5658667292b74ca19295827c2b22a0e903db283998ae2>
<https://camo.githubusercontent.com/3d9666a8f0c5658667292b74ca19295827c2b22a0e903db283998ae2>

```
In [68]: input_image = cv2.imread('MNIST_digit.png')
```

```
In [69]: type(input_image)
```

Out[69]: numpy.ndarray

```
In [70]: print(input_image)
```

```
[[[0 0 0]
  [0 0 0]
  [0 0 0]
  ...
  [0 0 0]
  [0 0 0]
  [0 0 0]]

 [[0 0 0]
  [0 0 0]
  [0 0 0]
  ...
  [0 0 0]
  [0 0 0]
  [0 0 0]]

 [[0 0 0]
  [0 0 0]
  [0 0 0]
  ...
  [0 0 0]
  [0 0 0]
  [0 0 0]]


 ...

 [[0 0 0]
  [0 0 0]
  [0 0 0]
  ...
  [0 0 0]
  [0 0 0]
  [0 0 0]]

 [[0 0 0]
  [0 0 0]
  [0 0 0]
  ...
  [0 0 0]
  [0 0 0]
  [0 0 0]]

 [[0 0 0]
  [0 0 0]
  [0 0 0]
  ...
  [0 0 0]
  [0 0 0]
  [0 0 0]]]
```

```
In [87]: plt.imshow(input_image)
```



```
Out[87]: <matplotlib.image.AxesImage at 0x25a1a2c1d50>
```

```
In [74]: input_image.shape
```

```
Out[74]: (318, 318, 3)
```

```
In [75]: grayscale = cv2.cvtColor(input_image, cv2.COLOR_RGB2GRAY)
```

```
In [76]: grayscale.shape
```


```
Out[76]: (318, 318)
```

```
In [77]: input_image_resize = cv2.resize(grayscale, (28, 28))
```

```
In [78]: input_image_resize.shape
```

```
Out[78]: (28, 28)
```

```
In [88]: plt.imshow(input_image_resize)
```



```
Out[88]: <matplotlib.image.AxesImage at 0x25a1a3534d0>
```

```
In [81]: input_image_resize = input_image_resize/255
```

```
In [82]: type(input_image_resize)
```

```
Out[82]: numpy.ndarray
```

```
In [83]: image_resized = np.reshape(input_image_resize, [1,28,28])
```

```
In [84]: input_prediction = model.predict(image_resized)
print(input_prediction)
```

```
1/1 [=====] - 0s 266ms/step
[[1.4969674e-17 6.4147948e-10 3.2105504e-06 1.0000000e+00 2.6870534e-15
 3.6363667e-06 5.8934504e-23 1.8358632e-04 5.4213418e-05 3.0160829e-01]]
```

```
In [85]: input_pred_label = np.argmax(input_prediction)
```

```
In [86]: print(input_pred_label)
```

3

Predictive System

```
In [ ]: input_image_path = input('Path of the image to be predicted: ')

input_image = cv2.imread(input_image_path)

cv2.imshow(input_image)

grayscale = cv2.cvtColor(input_image, cv2.COLOR_RGB2GRAY)

input_image_resize = cv2.resize(grayscale, (28, 28))

input_image_resize = input_image_resize/255


image_resized = np.reshape(input_image_resize, [1,28,28])

input_prediction = model.predict(image_resized)

input_pred_label = np.argmax(input_prediction)

print('The Handwritten Digit is recognised as ', input_pred_label)
```

Path of the image to be predicted: /content/MNIST_digit.png



The Handwritten Digit is recognised as 3

In []: